



Petabyte Scale Data Warehousing Greenplum

PLContainer --Trusting Untrusted Languages

Postgres Conf 2018

Marshall Presser

Craig Sylvester

Andreas Scherbaum

17 April 2018

Agenda

- The Problem
- What is PL/Container
- How to use PL/Container
- PL/Container Internals
- Future Work
- Q+A



**Pivotal
Greenplum Database**



PL/Python

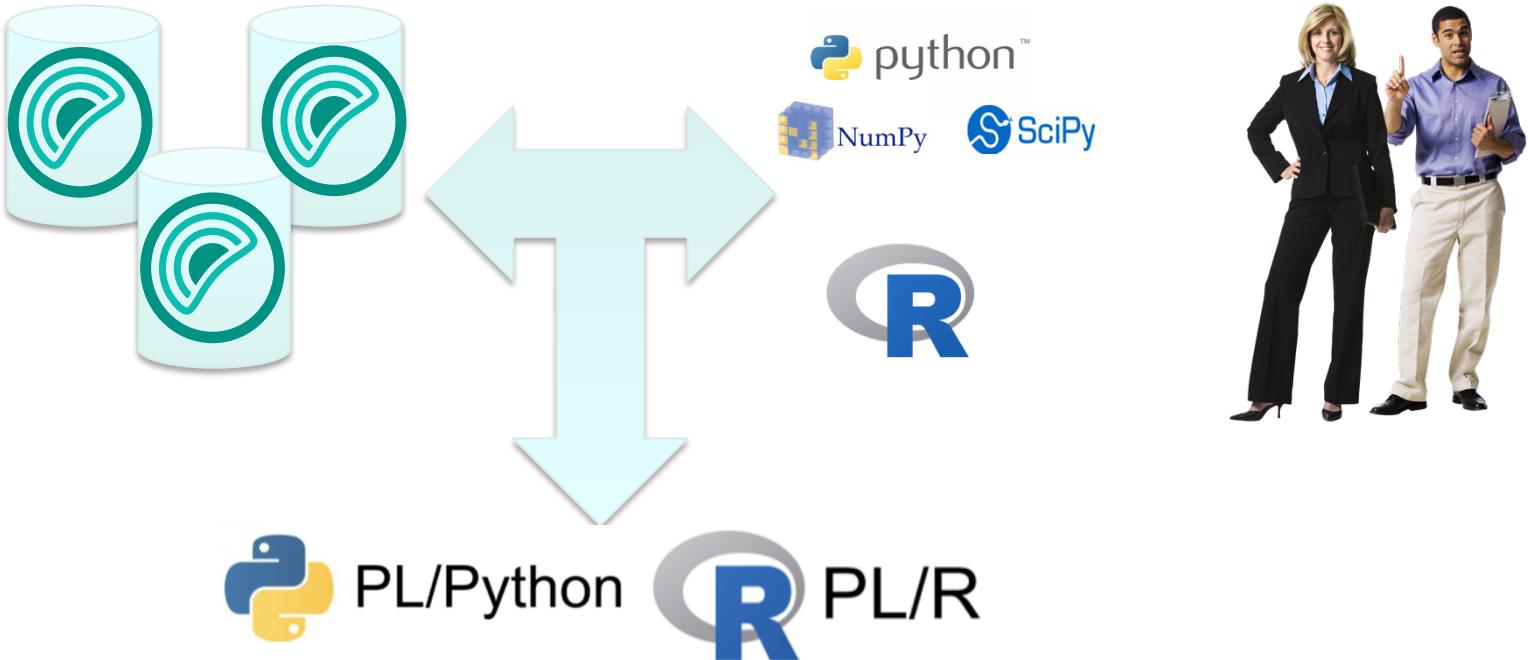


PL/R



PL/Container

We want to analyse data IN Database



BUT ...

PL/Python and PL/R are UNTRUSTED Languages

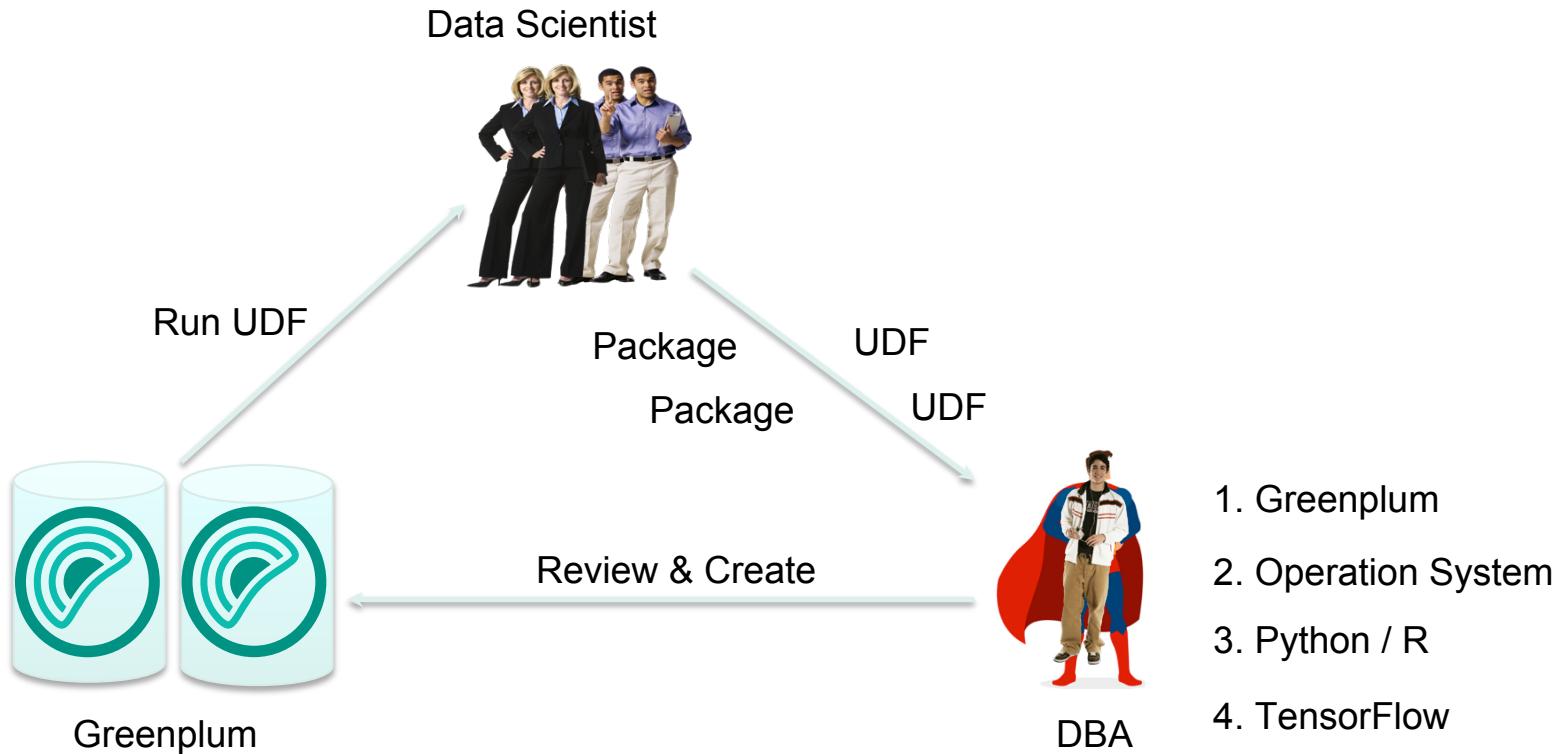
Only Superuser can Create UDF in Untrusted Languages



System("rm -rf /data")



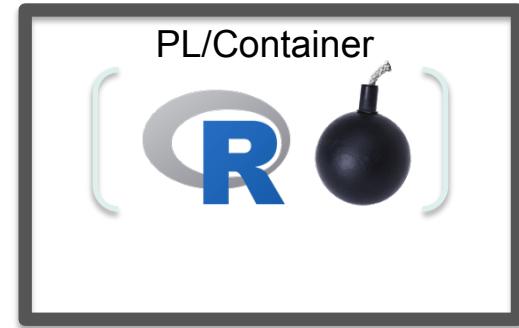
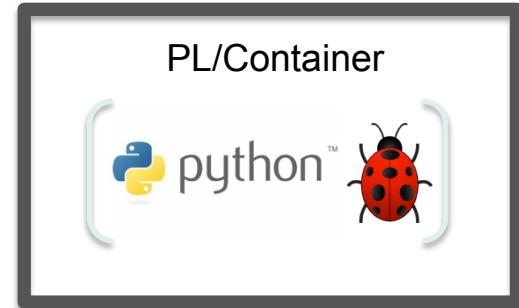
The Problem: Triangle Dependency



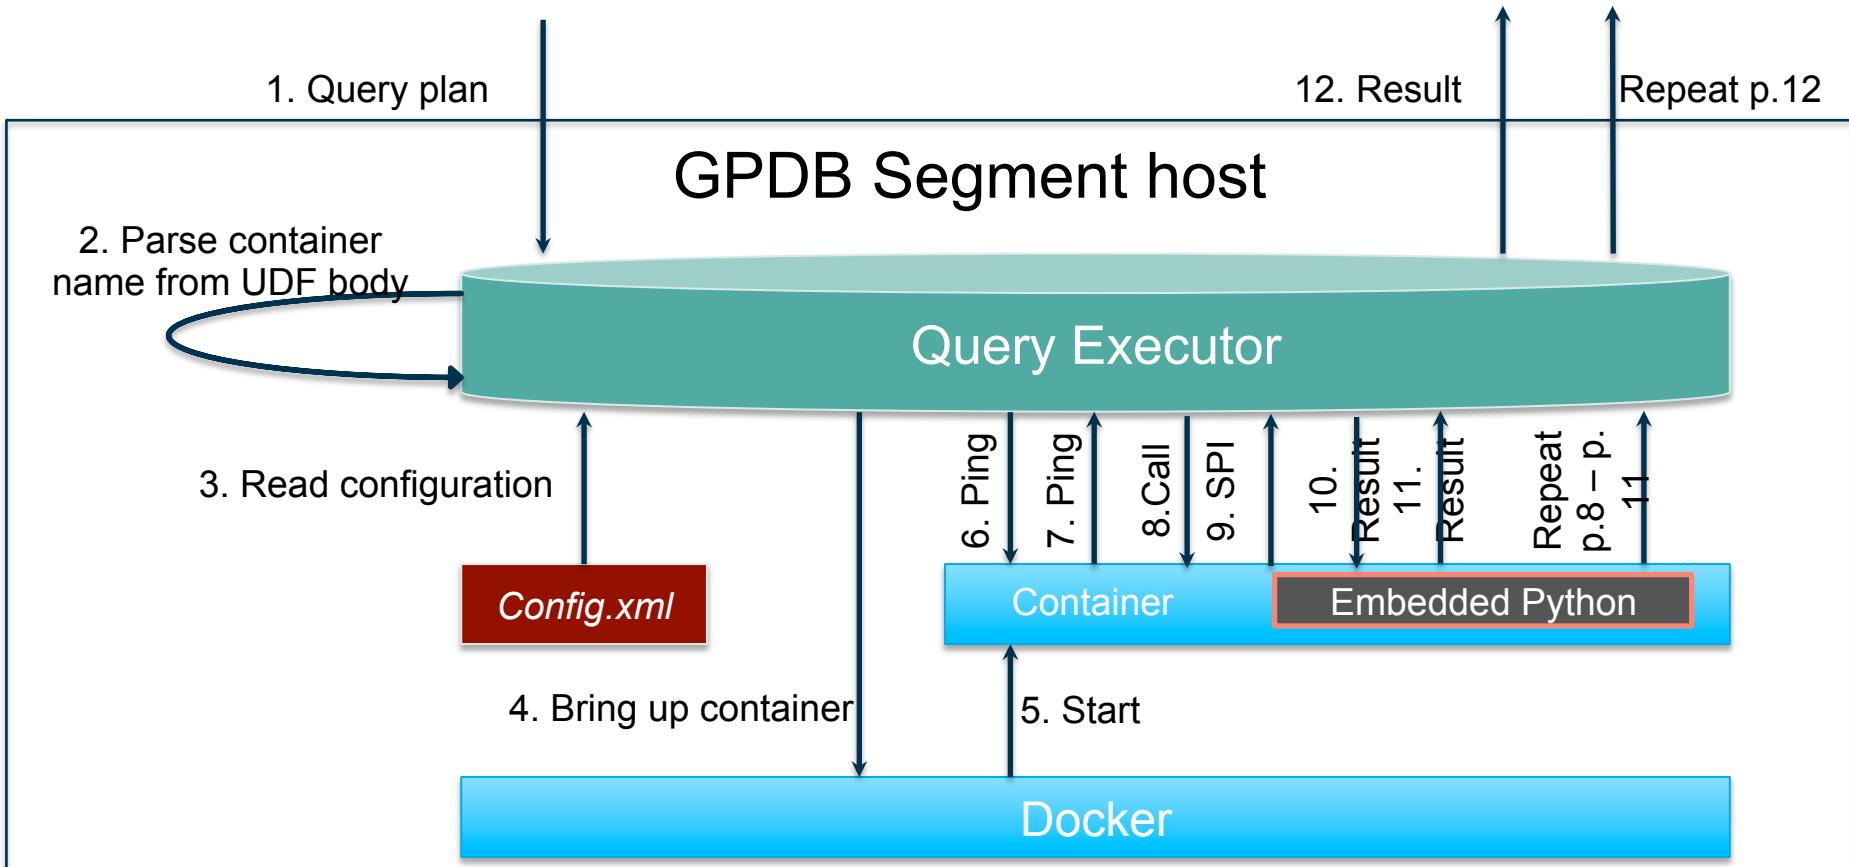
What is PL/Container?

PL/Container is a customizable, secure runtime for Greenplum Database Procedural Languages.

- Greenplum Database Extension
- Stateless
- Based on Docker Container
- Customizable
- Secure
- Isolated



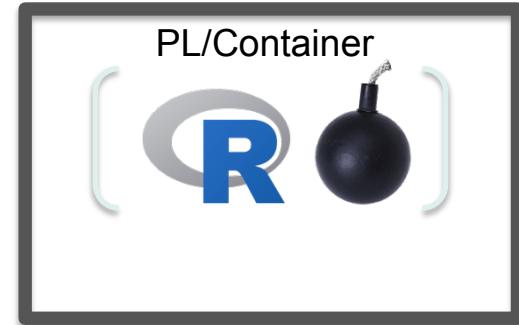
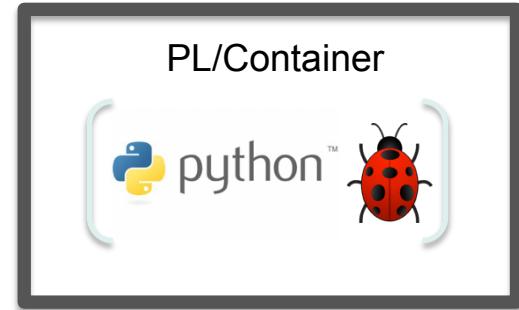
PL/Container Architecture



How UDF run in PL/Container?

PL/Container is a customizable, secure runtime for Greenplum Database Procedural Languages.

- a. PL/Container Extension starts a docker container (only in 1st call)
- b. Transfer UDF and data to docker container
- c. Run the UDF in docker container
- d. Contact the docker container to get the results



Install PL/Container on Greenplum

Install from Source Code

- source \$GPHOME/greenplum_path.sh
- make install

Install from GPPKG

- gppkg -i plcontianer-1.1.0-rhel7-x86_64.gppkg
- no additional dependencies



Prerequisites

Platform

Centos 6.6+ or 7.x

Database

Greenplum 5.2+

Docker

Docker 17.05+ on Centos7

Docker 1.7+ on Centos6

Build Custom Docker Image (optional)

Minimum Requirement:

- Python or R environment
- Add location of libpython.so and libR.so to LD_LIBRARY_PATH

```
FROM continuumio/anaconda
```

```
ENV LD_LIBRARY_PATH "/opt/conda/lib:$LD_LIBRARY_PATH"
```

Customize Your image:

- Install specific packages

```
FROM continuumio/anaconda3
```

```
RUN conda install -c conda-forge -y tensorflow
```

```
ENV LD_LIBRARY_PATH "/opt/conda/lib:/usr/local/lib:$LD_LIBRARY_PATH"
```

Configure PL/Container



XML

runtime

<id>

<image>

<command>

<shared directory>

<setting memory_mb>

<setting cpu_share>

container cgroup node

memory.memsw.limit_in_bytes

cpu.shares



Image

image add

image delete

image list



Runtime

runtime add

runtime delete

runtime backup

runtime restore

runtime edit

runtime show

Run PL/Container

Running a simple plpython UDF to calculate \log_{10}

```
postgres=# CREATE LANGUAGE plpythonu; plcontainer;

postgres=# CREATE OR REPLACE FUNCTION pylog10(input double precision) RETURNS double
precision AS $$

import math

return math.log10(input)

$$ LANGUAGE plpythonu;

postgres=# SELECT pylog10(100);
      pylog10
-----
              2
(1 row)
```

Run PL/Container

Running a simple PL/Container UDF to calculate \log_{10}

```
postgres=# CREATE EXTENSION plcontainer;

postgres=# CREATE OR REPLACE FUNCTION pylog10(Input double precision) RETURNS double
precision AS $$

# container: plc_python_shared

import math

return math.log10(input)

$$ LANGUAGE plcontainer;

postgres=# SELECT pylog10(100);
pylog10
-----
2
(1 row)
```

Performance

Test Environment

- Hardware: 6 virtual machines, each with 19G memory and 5 processors. (Intel(R) Xeon(R) CPU E5-2697 v2 @ 2.70GHz)
- Software: Centos7, GPDB 5.2 with 30 segments.

Workloads

- Long-running function
- Large input array function
- Large output array function

Performance

Long-running function

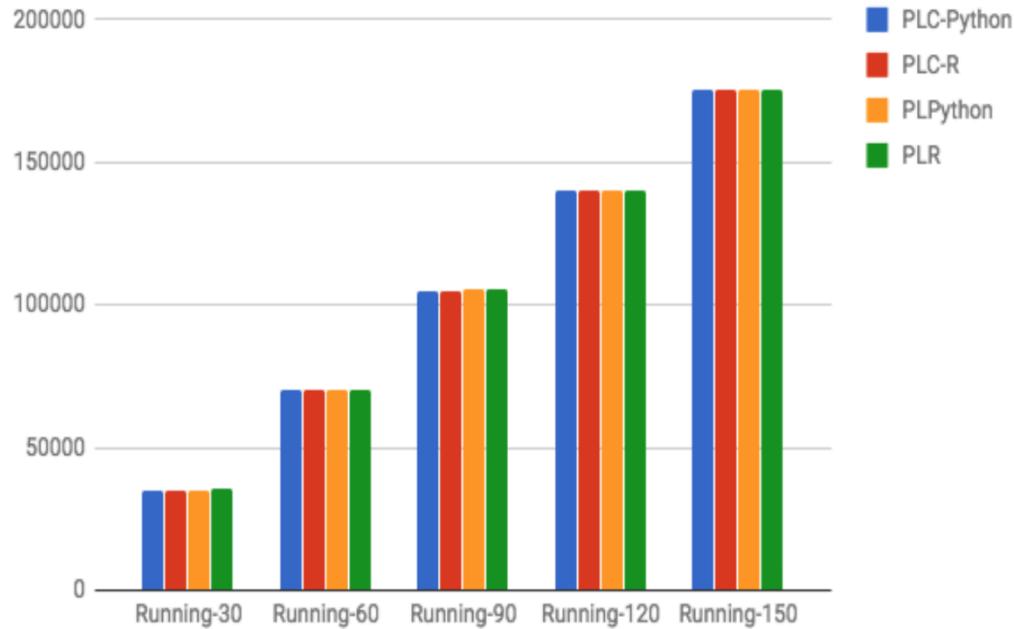
```
CREATE OR REPLACE FUNCTION pysleep(i int)
RETURNS void AS $$

# contaziner: plc_python_shared
import time
time.sleep(i)

$$ LANGUAGE plcontainer;
```

```
SELECT count(pysleep(1)) FROM tbl;
```

Long-time Running-Segment



no performance downgrade



Pivotal®

Transforming How The World Builds Software